MA783 Advanced Stochastic Processes Notes

Instructor: Prof. Solesne Bourguin

Student: Wancheng Lin

Contents

---

1 Probability Spaces, Random Variables, and Stochastic Processes

Sigma-Algebras and Probability Measures

Definition 1.1 (σ-algebra)**.** *Let Ω be a non-empty set. A family $\mathcal{F}$ of subsets of Ω is a* σ-algebra *if:*

1. $\Omega \in \mathcal{F}$, $\emptyset \in \mathcal{F}$,
2. $A \in \mathcal{F} \implies A^c := \Omega \setminus A \in \mathcal{F}$,
3. *If* $A_1, A_2, \dots \in \mathcal{F}$, *then* $\bigcup_{i=1}^{\infty} A_i \in \mathcal{F}$.

This structure guarantees closure under complements and countable unions, providing the minimal requirement for probability to be well-defined.

Definition 1.2 (Probability Space)**.** *A* probability space *is a triple $(\Omega, \mathcal{F}, P)$ where*

- $(\Omega, \mathcal{F})$ *is a measurable space,*
- $P : \mathcal{F} \to [0,1]$ *is a probability measure such that*

$$P(\Omega) = 1, \quad P\left(\bigcup_{i=1}^{\infty} A_i\right) = \sum_{i=1}^{\infty} P(A_i)$$

for all disjoint $A_i \in \mathcal{F}$.

A complete probability space also contains all subsets of null sets, ensuring that negligible events are measurable. This prevents technical issues later when working with almost sure properties.

Random Variables and Distributions

Definition 1.3 (Random Variable)**.** *A mapping $X : \Omega \to \mathbb{R}^n$ is a random variable if it is measurable, i.e.*

$$X^{-1}(B) := \{\omega \in \Omega : X(\omega) \in B\} \in \mathcal{F}, \quad \forall B \in \mathcal{B}(\mathbb{R}^n),$$

where $\mathcal{B}(\mathbb{R}^n)$ is the Borel σ-algebra.

This guarantees that events of the form $\{X \in B\}$ are measurable. The σ-algebra generated by X, written $\sigma(X)$, represents the information revealed by knowing X.

Definition 1.4 (Distribution)**.** *The distribution of X is the probability measure μ_X on $\mathcal{B}(\mathbb{R}^n)$ defined by*

$$\mu_X(B) := P\big(X^{-1}(B)\big), \quad B \in \mathcal{B}(\mathbb{R}^n).$$

Thus the distribution is the pushforward measure of P under X.

Definition 1.5 (Expectation)**.** *If $\int_\Omega |X(\omega)|\, dP(\omega) < \infty$, the expectation of X is*

$$E[X] = \int_\Omega X(\omega)\, dP(\omega) = \int_{\mathbb{R}^n} x\, d\mu_X(x).$$

L^p Spaces and Equivalence Classes

For $p \geq 1$, the L^p-norm of a random variable X is defined by

$$\|X\|_p := \Big(\int_\Omega |X(\omega)|^p\, dP(\omega) \Big)^{1/p}.$$

Definition 1.6 (L^p Spaces)**.** *The space $L^p(P)$ consists of all random variables $X : \Omega \to \mathbb{R}$ with $\|X\|_p < \infty$.*

Equivalence classes. Strictly speaking, $L^p(P)$ is not the set of all random variables with finite p-moment, but rather the set of their *equivalence classes*. Two random variables X and Y are identified whenever

$$P(X = Y) = 1,$$

i.e. they differ only on a null set. Thus, each element of $L^p(P)$ is an equivalence class

$$[X] := \{Y : P(X = Y) = 1\}.$$

This convention avoids ambiguity: if X and Y agree almost surely, then all probabilistic properties we care about (expectation, variance, stochastic integrals) are the same. Working with equivalence classes ensures that the L^p-norm is a true norm rather than merely a seminorm.

Banach and Hilbert structures. The space $L^p(P)$ with norm $\|\cdot\|_p$ is complete, hence a Banach space. When $p = 2$, the space becomes a Hilbert space with inner product

$$\langle X, Y \rangle := E[XY],$$

well-defined on equivalence classes. In particular,

$$\|X\|_2^2 = \langle X, X \rangle.$$

Why L^2 matters in stochastic calculus. Itô integrals are first defined for elementary processes and then extended to all integrands in L^2. The Hilbert structure is essential: convergence in L^2 allows us to use projection and orthogonality arguments. Moreover, martingale theory and isometries (e.g. Itô's isometry)

$$E\left[\left(\int_0^T f(t)\, dB_t\right)^2\right] = E\left[\int_0^T f(t)^2\, dt\right]$$

rely explicitly on the L^2 framework.

Interpretation. From a probabilistic viewpoint, the use of equivalence classes formalizes the phrase "almost surely." We are not interested in the exact values of random variables on exceptional sets of probability zero, since they do not affect expectations, variances, or distributions. Thus L^p spaces provide a natural analytic setting for stochastic processes.

Stochastic Processes

A *stochastic process* $\{X_t\}_{t \in T}$ is simply a family of random variables, all defined on the same probability space $(\Omega, \mathcal{F}, P)$ and taking values in $\mathbb{R}^n$.

Two complementary perspectives. There are always two natural ways to look at such a process:

- If we fix t, then X_t is just a random variable. We can study its distribution, compute its expectation, variance, etc.
- If we fix $\omega \in \Omega$, then $t \mapsto X_t(\omega)$ is a deterministic function of t, called the sample path of the process for that outcome ω.

This duality — randomness across ω and dynamics across t — is what makes stochastic processes both rich and subtle.

Intuitive picture. It is often helpful to think of t as representing *time* and ω as representing an individual "particle" or "experiment." Then $X_t(\omega)$ is the state of that particle at time t. Sometimes one also writes $X(t, \omega)$ instead of $X_t(\omega)$, emphasizing that the process is really a function

$$(t, \omega) \mapsto X(t, \omega),$$

defined on $T \times \Omega$ and taking values in $\mathbb{R}^n$. This viewpoint will become important later, since joint measurability in (t, ω) is crucial for defining stochastic integrals.

Path space representation. Fixing ω, we can identify it with the entire path $t \mapsto X_t(\omega)$. In this sense, we may think of Ω as a subset of the path space

$$(\mathbb{R}^n)^T = \{\text{all functions } f : T \to \mathbb{R}^n\}.$$

The σ-algebra $\mathcal{F}$ will then contain the cylinder σ-algebra $\mathcal{B}$ generated by sets of the form

$$\{\omega : \omega(t_1) \in F_1, \ldots, \omega(t_k) \in F_k\}, \quad F_i \in \mathcal{B}(\mathbb{R}^n).$$

From this point of view, a stochastic process is nothing but a probability measure P on the measurable space $\big((\mathbb{R}^n)^T, \mathcal{B}\big)$.

Finite-dimensional distributions. The family of distributions of the vectors $(X_{t_1}, \dots, X_{t_k})$, for $t_1, \dots, t_k \in T$, are called the *finite-dimensional distributions* of the process. Explicitly,

$$\mu_{t_1,\dots,t_k}(F_1 \times \cdots \times F_k) = P(X_{t_1} \in F_1, \dots, X_{t_k} \in F_k)\,,$$

for Borel sets $F_i \subseteq \mathbb{R}^n$. These distributions capture many important properties of the process — for instance, whether increments are Gaussian or independent — though not everything (they do not by themselves encode path continuity, for example).

Kolmogorov's extension theorem. Conversely, if we are given a family of candidate finite-dimensional distributions $\{\nu_{t_1,\dots,t_k}\}$, it is natural to ask: does there exist a stochastic process that realizes them? Kolmogorov's extension theorem gives a celebrated answer: if the family satisfies two natural consistency conditions (marginalization and permutation symmetry), then there exists a stochastic process $\{Y_t\}_{t\in T}$ whose finite-dimensional distributions coincide with $\{\nu_{t_1,\dots,t_k}\}$.

This result is fundamental: it allows us to construct processes like Brownian motion purely from the specification of their finite-dimensional laws, before we even worry about pathwise properties like continuity. For details and proofs, see Lamperti (1977) or Kallenberg (2002).

Independence

Two events $A, B \in \mathcal{F}$ are independent if

$$P(A \cap B) = P(A)P(B).$$

Two random variables X, Y are independent if the σ-algebras $\sigma(X)$ and $\sigma(Y)$ are independent. For independent X and Y, provided $E[|X|], E[|Y|] < \infty$, we have

$$E[XY] = E[X]\,E[Y].$$

Conditional Expectation

Let $H \subseteq \mathcal{F}$ be a sub-σ-algebra. For $X \in L^1(P)$, the conditional expectation $E[X|H]$ is the H-measurable random variable satisfying

$$\int_H E[X|H]\,dP = \int_H X\,dP, \quad \forall H \in H.$$

It is unique up to null sets and inherits key properties:

$$\begin{aligned} E[aX + bY \mid H] &= aE[X|H] + bE[Y|H], \\ E[E[X|H]] &= E[X], \\ E[X|H] &= X \quad \text{if } X \text{ is } H\text{-measurable}, \\ E[X|H] &= E[X] \quad \text{if } X \text{ independent of } H, \\ E[XY|H] &= Y\,E[X|H] \quad \text{if } Y \text{ is } H\text{-measurable}. \end{aligned}$$

Moreover, Jensen's inequality extends to this setting:

$$\varphi(E[X|H]) \le E[\varphi(X)|H], \quad \text{for convex } \varphi.$$

Filtrations and Adapted Processes

A *filtration* is an increasing family of sub-σ-algebras $\{\mathcal{F}_t\}_{t\ge 0}$ representing the information available up to time t.

A process $\{X_t\}$ is said to be adapted if X_t is $\mathcal{F}_t$-measurable for all t. This notion formalizes the idea that the process does not "look into the future."

For example, if $\{B_t\}$ is Brownian motion, then $\mathcal{F}_t^B = \sigma(B_s : 0 \le s \le t)$ is the natural filtration of the process, representing exactly the history of the motion up to time t.

2 Brownian Motion and Itô Integration

Construction of Brownian Motion

To construct a Brownian motion $\{B_t : t \geq 0\}$ using Kolmogorov's extension theorem, we must specify a consistent family of finite-dimensional distributions $\{\nu_{t_1,\dots,t_k}\}$.

For $x, y \in \mathbb{R}^n$ and $t > 0$, define the Gaussian transition kernel

$$p(t,x,y) = (2\pi t)^{-n/2} \exp\left(-\frac{|x-y|^2}{2t}\right), \qquad p(0,x,y) = \delta_x(y).$$

For $0 \leq t_1 \leq t_2 \leq \cdots \leq t_k$, define

$$\nu_{t_1,\dots,t_k}(F_1 \times \cdots \times F_k) = \int_{F_1 \times \cdots \times F_k} p(t_1,x,x_1)p(t_2-t_1,x_1,x_2)\cdots p(t_k-t_{k-1},x_{k-1},x_k)\,dx_1\cdots dx_k.$$

By permutation symmetry, this definition extends to any ordering of times and automatically satisfies the first consistency condition.

Remark. The second consistency condition is also satisfied: marginalizing out later coordinates leaves the earlier distribution unchanged, thanks to $\int_{\mathbb{R}^n} p(t,x,y)dy = 1$.

Theorem 2.1 (Kolmogorov Extension)**.** *There exists a probability space $(\Omega, \mathcal{F}, P^x)$ and a stochastic process $\{B_t : t \geq 0\}$ such that the finite-dimensional distributions of B_t coincide with $\nu_{t_1,\dots,t_k}$ constructed above. Moreover $P^x(B_0 = x) = 1$.*

Definition 2.1. *Any such process is called a (version of) n-dimensional Brownian motion starting at x.*

Remark. This construction does not produce a unique process — versions may differ in pathwise properties. We choose a version with continuous paths, justified by Kolmogorov's continuity theorem.

Kolmogorov's Continuity Theorem

Theorem 2.2 (Kolmogorov's Continuity)**.** *Suppose $\{X_t\}_{t\geq 0}$ is a process such that for all $T > 0$ there exist $\alpha, \beta, D > 0$ with*

$$E\big(|X_t - X_s|^\alpha\big) \leq D|t-s|^{1+\beta}, \quad 0 \leq s,t \leq T.$$

Then X admits a continuous version.

For Brownian motion,

$$E^x\big(|B_t - B_s|^4\big) = n(n+2)|t-s|^2,$$

so the theorem holds with $\alpha = 4$, $\beta = 1$, $D = n(n+2)$. Thus Brownian motion always has a continuous version.

Remark. If $B_t = (B_t^{(1)}, \dots, B_t^{(n)})$ is n-dimensional Brownian motion, then the coordinate processes $\{B_t^{(i)}\}$ are independent one-dimensional Brownian motions.

Basic Properties of Brownian Motion

Brownian motion $\{B_t\}_{t\geq 0}$ is defined by its finite-dimensional distributions. From this definition we can immediately deduce a number of fundamental properties.

1. Gaussian process. Brownian motion is a Gaussian process. This means that for every finite set of time points

$$0 \leq t_1 \leq \cdots \leq t_k,$$

the random vector

$$Z = (B_{t_1}, \dots, B_{t_k}) \in \mathbb{R}^{nk}$$

is multivariate normal.

The mean vector and covariance matrix are given by

$$M = \big(E[B_{t_1}], \dots, E[B_{t_k}]\big), \qquad C = [c_{jm}], \quad c_{jm} = E\big[(B_{t_j} - M_j)(B_{t_m} - M_m)\big].$$

Equivalently, the joint characteristic function has the familiar Gaussian form:

$$E^x\left[\exp\left(i\sum_{j=1}^{nk} u_j Z_j\right)\right] = \exp\Big(iu^\top M - \tfrac{1}{2}u^\top C u\Big).$$

In particular, for one-dimensional Brownian motion starting at x,

$$E^x[B_t] = x, \qquad \mathrm{Var}^x(B_t) = E^x\big[(B_t - x)^2\big] = t.$$

In the n-dimensional case, $\mathrm{Var}^x(B_t) = nt$.

2. Covariance structure. For $0 \le s \le t$, the covariance of increments is

$$E^x\big[(B_t - B_s)^2\big] = n(t-s).$$

More generally, for $i, j \in \{1, \dots, n\}$,

$$\mathrm{Cov}(B_t^{(i)}, B_s^{(j)}) = \delta_{ij} \min(s,t),$$

where δ_{ij} is the Kronecker delta. This formula reflects two key features: different coordinates are independent, and covariance grows linearly with the overlap in time.

3. Independent increments. One of the defining properties of Brownian motion is that increments over disjoint intervals are independent. Specifically, for

$$0 \le t_1 < t_2 < \cdots < t_k,$$

the random variables

$$B_{t_1}, \quad B_{t_2} - B_{t_1}, \quad \dots, \quad B_{t_k} - B_{t_{k-1}}$$

are independent.

4. Stationary Gaussian increments. Not only are increments independent, they are also stationary and Gaussian. For $0 \le s < t$,

$$B_t - B_s \sim \mathcal{N}(0, (t-s)I_n).$$

Thus the distribution of an increment depends only on the length of the interval, not on its location in time.

Summary. Brownian motion is therefore:

- a Gaussian process with mean x and covariance $\min(s,t)I_n$,
- with continuous paths (after choosing the continuous version),
- and with stationary, independent, normally distributed increments.

These properties uniquely characterize Brownian motion and underlie its role as the canonical model of continuous-time noise.

Remark. If $B_t = \big(B_t^{(1)}, \dots, B_t^{(n)}\big)$ is an n-dimensional Brownian motion, then each coordinate process $\{B_t^{(j)}\}_{t \ge 0}$, $1 \le j \le n$, is itself a one-dimensional Brownian motion. Moreover, these coordinate processes are mutually independent. This follows directly from the covariance structure, since

$$\mathrm{Cov}(B_t^{(i)}, B_s^{(j)}) = \delta_{ij} \min(s,t).$$

Continuity and Versions of Processes

Definition 2.2 (Version / Modification)**.** *Let $\{X_t\}_{t \ge 0}$ and $\{Y_t\}_{t \ge 0}$ be stochastic processes defined on the same probability space $(\Omega, \mathcal{F}, P)$. We say that $\{X_t\}$ is a* version *(or* modification*) of $\{Y_t\}$ if for every fixed $t \ge 0$,*

$$P\big(\{\omega \in \Omega : X_t(\omega) = Y_t(\omega)\}\big) = 1.$$

Thus, versions agree almost surely at each fixed time t, although they may differ on a null set that can vary with t. Importantly, versions have the same finite-dimensional distributions (f.d.d.'s).

Finite-dimensional distributions. Given a process $\{X_t\}_{t\geq 0}$, the collection of joint distributions of

$$(X_{t_1}, \ldots, X_{t_k}) \quad \text{for all } k \geq 1,\ 0 \leq t_1 < \cdots < t_k$$

is called the system of finite-dimensional distributions. If $\{X_t\}$ and $\{Y_t\}$ are versions, then clearly

$$(X_{t_1}, \ldots, X_{t_k}) \overset{d}{=} (Y_{t_1}, \ldots, Y_{t_k}), \quad \forall k, t_1, \ldots, t_k,$$

so they induce the same probability law on $\mathbb{R}^k$.

Why versions matter. The definition of Brownian motion via Kolmogorov's extension theorem yields only a process with the correct f.d.d.'s. However, this construction does not guarantee any regularity of sample paths: the raw object might be highly irregular, even discontinuous almost everywhere. To obtain the "classical" Brownian motion with continuous paths, one needs to show that there exists a version of the process that is continuous with probability one. This is where Kolmogorov's continuity theorem enters.

Theorem 2.3 (Kolmogorov's Continuity Theorem)**.** *Let $\{X_t : t \in T\}$ be a stochastic process such that for some $\alpha, \beta, D > 0$,*

$$\mathbb{E}\big[|X_t - X_s|^\alpha\big] \ \leq\ D\,|t-s|^{1+\beta}, \quad \forall s, t \in T.$$

Then there exists a continuous version $\{\tilde{X}_t\}$ of $\{X_t\}$ such that

$$P\big(X_t = \tilde{X}_t,\ \forall t \in T\big) = 1.$$

Moreover, the paths of $\tilde{X}_t$ are Hölder continuous of any order $\gamma < \frac{\beta}{\alpha}$.

Application to Brownian motion. For Brownian motion $B_t \in \mathbb{R}^n$, one computes

$$\mathbb{E}\big[|B_t - B_s|^4\big] = n(n+2)\,|t-s|^2.$$

Thus the continuity theorem applies with $\alpha = 4$, $\beta = 1$, and $D = n(n+2)$. Hence Brownian motion admits a continuous version with paths that are almost surely Hölder continuous of any exponent $\gamma < \frac{1}{4}$.

Key insight. Versions emphasize the distinction between:

- the *law* of a process, determined by its finite-dimensional distributions, and
- the *sample path properties*, such as continuity, differentiability, or bounded variation.

The Kolmogorov theorem is indispensable: it elevates the abstract Brownian motion from a mere collection of distributions to the concrete object with continuous trajectories that we use in stochastic calculus.

In particular, when we speak of "Brownian motion" in stochastic analysis, we always mean this continuous version.

Some Reflection on Canonical Brownian Motion

The Brownian motion defined via Kolmogorov's extension theorem is not unique. Indeed, there may exist several probability spaces $(\Omega, \mathcal{F}, P^x)$ and processes $\{B_t\}$ such that the finite-dimensional distributions satisfy the conditions of Brownian motion. However, for our purposes this non-uniqueness is not problematic: we may simply choose a convenient version to work with.

Continuous paths and identification. As established by Kolmogorov's continuity theorem, there exists a version of Brownian motion with continuous paths almost surely. Thus, for almost all $\omega \in \Omega$, we may identify ω with a continuous function

$$t \mapsto B_t(\omega), \quad t \in [0, \infty),\ B_t(\omega) \in \mathbb{R}^n.$$

Hence, we may regard Brownian motion as a probability measure P^x on the path space $C([0,\infty), \mathbb{R}^n)$. This version is called the *canonical Brownian motion*.

Why canonical? This point of view is not only intuitive but also technically advantageous. The space $C([0,\infty),\mathbb{R}^n)$ of continuous functions, equipped with the topology of uniform convergence on compact sets, is a Polish space (complete and separable metric space). This allows us to employ powerful results from measure theory and probability on Polish spaces, and is the starting point for much of the modern theory of stochastic processes (see Stroock and Varadhan (1979)).

Subtlety of measurability. At first glance, one might ask whether

$$t \mapsto B_t(\omega) \quad \text{is continuous for almost all } \omega.$$

However, the set

$$H = \{\omega \in \Omega : t \mapsto B_t(\omega) \text{ is continuous}\}$$

is not measurable with respect to the canonical product σ-algebra $\mathcal{B}(\mathbb{R}^n)^{[0,\infty)}$, since it involves uncountably many time indices.

By reformulating the construction in the canonical path space $C([0,\infty),\mathbb{R}^n)$, this measurability issue disappears: continuity is built into the path space itself. Thus, the canonical construction provides a rigorous and convenient framework in which Brownian motion is viewed as a random element of $C([0,\infty),\mathbb{R}^n)$.

From Discrete Models to Stochastic Differential Equations

We began with a deterministic growth model

$$\frac{dN(t)}{dt} = r(t)\,N(t),$$

which describes the rate of change of a population $N(t)$ with deterministic growth rate $r(t)$.

In realistic settings, however, the growth rate is subject to random environmental fluctuations. A natural idea is to add a "noise term":

$$\frac{dN(t)}{dt} = \big(r(t) + \text{noise}\big)\,N(t).$$

Generalization. For a general state process $\{X_t\}$, we may write

$$\frac{dX_t}{dt} = b(t, X_t) + \sigma(t, X_t) \cdot \text{noise},$$

where

- $b(t, x)$ is the *drift* (deterministic trend),
- $\sigma(t, x)$ scales the *random fluctuations*.

Noise as Brownian motion. To make this precise, we need to model "noise" as a well-defined stochastic process. The canonical choice is one-dimensional Brownian motion $\{W_t\}$, characterized by:

1. $W_0 = 0$ almost surely,
2. Independent increments: $W_{t_2} - W_{t_1}$ is independent of the past if $t_2 > t_1$,
3. Stationary increments: $W_{t+s} - W_s \sim \mathcal{N}(0, t)$ for all $s, t \geq 0$,
4. $E[W_t] = 0$ and $\text{Var}(W_t) = t$.

Thus the natural stochastic differential equation (SDE) is

$$dX_t = b(t, X_t)\,dt + \sigma(t, X_t)\,dW_t.$$

This is the standard Brownian-driven model, which forms the foundation for Itô calculus.

Motivation for the Itô Integral

To interpret an SDE rigorously, we seek processes $\{X_t\}$ satisfying

$$X_t = X_0 + \int_0^t b(s, X_s)\,ds + \int_0^t \sigma(s, X_s)\,dB_s.$$

The first integral is classical (Riemann or Lebesgue), but the second is new: a stochastic integral with respect to Brownian motion.

The central problem. Brownian paths are almost surely continuous but nowhere differentiable and of infinite variation. Therefore, the last term cannot be understood as an ordinary Riemann–Stieltjes integral. We must instead construct a new theory of integration that is compatible with the quadratic variation of Brownian motion:

$$[B]_t = t.$$

Interpretation. The stochastic term

$$\int_0^t \sigma(s, X_s)\,dB_s$$

represents the accumulated effect of "white noise" fluctuations scaled by σ. This captures how randomness enters continuously into the dynamics of X_t.

Conclusion. The task of the next sections is to rigorously define the Itô integral, first for simple adapted processes and then extending to general square-integrable integrands. This construction will provide the foundation for the theory of SDEs.

Why the Riemann–Stieltjes Approach Fails

A natural first attempt to define the stochastic integral

$$\int_0^T f(t,\omega)\, dB_t(\omega)$$

is to view $t \mapsto B_t(\omega)$ as a continuous function for each fixed ω, and try to interpret this as a classical Riemann–Stieltjes (RS) integral of the form $\int f\, dg$.

Indeed, if g has bounded variation, the RS integral $\int f\, dg$ exists whenever f is continuous. However, Brownian paths have two pathological features:

1. They are *nowhere differentiable*, so the interpretation $\int f(t)\, dB_t = \int f(t)\, B_t' dt$ is meaningless.
2. They have *infinite total variation* almost surely:

$$V_{0,T}(B) = \sup_{\pi} \sum_{j=0}^{n-1} |B_{t_{j+1}} - B_{t_j}| = \infty \quad \text{a.s.},$$

   where the supremum runs over partitions π of $[0,T]$.

Because of (2), the RS integral cannot be applied. This failure is not only theoretical but can be demonstrated concretely by approximations.

Illustrative Example. Consider $f(t,\omega) = B_t(\omega)$. Define step-function approximations on dyadic intervals:

$$\phi_1(t,\omega) = \sum_{j\geq 0} B_{j2^{-n}}(\omega)\, 1_{[j2^{-n},(j+1)2^{-n})}(t),$$

$$\phi_2(t,\omega) = \sum_{j\geq 0} B_{(j+1)2^{-n}}(\omega)\, 1_{[j2^{-n},(j+1)2^{-n})}(t).$$

These correspond to left- and right-point Riemann sums.

Compute their integrals:

$$\int_0^T \phi_1(t,\omega)\, dB_t(\omega) = \sum_j B_{t_j}(\omega)\, (B_{t_{j+1}} - B_{t_j}),$$

$$\int_0^T \phi_2(t,\omega)\, dB_t(\omega) = \sum_j B_{t_{j+1}}(\omega)\, (B_{t_{j+1}} - B_{t_j}).$$

Taking expectations:

$$E\left[\int_0^T \phi_1\, dB\right] = \sum_j E[B_{t_j}(B_{t_{j+1}} - B_{t_j})] = 0,$$

since increments are independent of the past.

But

$$E\left[\int_0^T \phi_2\, dB\right] = \sum_j E[B_{t_{j+1}}(B_{t_{j+1}} - B_{t_j})] = \sum_j (t_{j+1} - t_j) = T.$$

Thus two reasonable RS-type approximations give different limits. This shows the RS definition is inconsistent for Brownian integrals.

Conceptual Reason. The breakdown arises because RS integrals are designed for integrators of bounded variation, while Brownian motion has variation "too large." In fact, instead of total variation, Brownian

motion has a well-defined *quadratic variation*:

$$[B]_t = \lim_{|\pi|\to 0} \sum_j (B_{t_{j+1}} - B_{t_j})^2 = t,$$

which suggests that an integration theory adapted to quadratic, not total, variation is required.

Conclusion. Therefore, the Riemann–Stieltjes framework is inadequate. The correct approach, pioneered by Itô, defines the integral by:

- restricting initially to *elementary adapted processes*,
- defining the integral via increments of Brownian motion,
- and extending by L^2-limits, using the quadratic variation structure.

This leads to the Itô integral, which is consistent, linear, and satisfies the crucial *Itô isometry*.

Preliminaries: Filtration and Adaptedness

Let $\{B_t\}$ be an n-dimensional Brownian motion. Define

$$\mathcal{F}_t = \sigma(B_s^{(i)} : 0 \le s \le t,\ 1 \le i \le n),$$

the natural filtration. Intuitively, $\mathcal{F}_t$ represents the history of the process up to time t. A process $\phi(t,\omega)$ is *adapted* if $\phi(t,\cdot)$ is $\mathcal{F}_t$-measurable for each t.

Admissible Integrands

We now specify the class of processes we may integrate.

Definition 2.3. *Let $V(S,T)$ denote the class of processes $f : [S,T] \times \Omega \to \mathbb{R}$ such that:*

1. *$(t,\omega) \mapsto f(t,\omega)$ is $\mathcal{B}([S,T]) \otimes \mathcal{F}$-measurable;*
2. *$f(t,\omega)$ is adapted to $\{\mathcal{F}_t\}$;*
3. *$E\left[\int_S^T f(t,\omega)^2\, dt\right] < \infty$.*

This ensures that f is square-integrable in time and adapted to the filtration.

Elementary Integrands and Definition

Definition 2.4 (Elementary Process)**.** *A function $\phi \in V(S,T)$ is elementary if it has the form*

$$\phi(t,\omega) = \sum_j e_j(\omega)\, 1_{(t_j, t_{j+1}]}(t),$$

where e_j is $\mathcal{F}_{t_j}$-measurable and square-integrable.

For such ϕ, define the stochastic integral by

$$\int_S^T \phi(t,\omega)\, dB_t(\omega) := \sum_j e_j(\omega)\, (B_{t_{j+1}} - B_{t_j}).$$

This definition parallels the Riemann sum approach, but crucially the coefficients are measurable with respect to the past, not the future.

Remark. If e_j depended on $B_{t_{j+1}}$, the integral would not be well-defined. Adaptedness ensures causality.

Itô Isometry

To extend the integral beyond elementary processes, we use the following fundamental identity.

Lemma 2.1 (Itô Isometry)**.** *If $\phi \in V(S,T)$ is elementary, then*

$$E\left[\left(\int_S^T \phi(t,\omega)\, dB_t\right)^2\right] = E\left[\int_S^T \phi(t,\omega)^2\, dt\right].$$

Sketch. For $\phi = \sum e_j 1_{(t_j,t_{j+1}]}$, note that

$$\int_S^T \phi(t)\, dB_t = \sum_j e_j (B_{t_{j+1}} - B_{t_j}).$$

Using independence and zero mean of increments:

$$E\big[e_i e_j (B_{t_{i+1}} - B_{t_i})(B_{t_{j+1}} - B_{t_j})\big] = 0, \quad i \ne j,$$

and for $i = j$,

$$E[e_j^2 (B_{t_{j+1}} - B_{t_j})^2] = E[e_j^2](t_{j+1} - t_j).$$

Summing yields the result. □

Extension to General Integrands

Using the isometry, we extend the definition by density.

[Itô Integral] For $f \in V(S,T)$, there exists a sequence of elementary processes ϕ_n with

$$E\left[\int_S^T (f(t) - \phi_n(t))^2 \, dt\right] \to 0.$$

Define

$$I(f) = \int_S^T f(t)\, dB_t := L^2 - \lim_{n\to\infty} \int_S^T \phi_n(t)\, dB_t.$$

The limit exists and is unique in $L^2(\Omega)$ by the Itô isometry.

Properties of the Itô Integral

- Linearity: $\int f + g \, dB = \int f \, dB + \int g \, dB$.
- Isometry: $E[(\int f \, dB)^2] = E[\int f^2 \, dt]$.
- Martingale property: $\int_0^t f(s)\, dB_s$ is a martingale.
- Zero mean: $E[\int f \, dB] = 0$.

Remark. This construction is robust: unlike Riemann–Stieltjes, the Itô integral accommodates the roughness of Brownian motion via square-integrability and martingale structure.